

```

1)# Question1.py
# import random

import numpy as np

def logisticregr(theta, n, m):
    """Function to generate  $n*(m+1)$  independent variables,  $n*1$  dependent
    variable, and coefficient matrix  $(m+1*1)$ 

    Args:
        theta (float): spread of noise in the output variable
        n (int): size of the data set
        m (int): number of independent variables

    Returns:
        array: independent, dependent variables matrix and coefficient
matrix
    """
    x_ones = np.ones((n, 1)) # first column of ones
    # np.random.seed(1998)
    x_random = np.random.randn(n, m) # (n*m) matrix of random Real numbers
    X = np.concatenate((x_ones, x_random), axis=1) # n rows and m+1
columns where 1st column is of ones
    beta = np.random.rand(m + 1, 1) # Coefficients
    Y = (
        1 / (1 + np.exp(-np.dot(X, beta))) >= 0.5
    ) # X*beta i.e. matrix multiplication of (n*m) and (m * 1) shape= Y
shape (n*1)

    Y = Y.astype(int)
    noise = np.random.binomial(n=1, p=theta, size=(n, 1))
    Y = Y + noise
    Y %= 2

    ##### ----- #####
    # Alternatively, we can use Maximum likelihood estimation to add noise
to Y

    # probflip = theta # probability that the output would be flipped

```

```

    # # As it is a bernoulli problem, by Maximum likelihood estimation, we
know (theta = # of flips/ n)
    # # hence, total number of flips = n*theta)...taking this number of
random sets of inputs whose output is reversed

    # # flipping outputs of Y (adding noise)
    # for i in random.sample(range(0, n), int(n * probflip)): # n unique
random intergers
    #     if Y[i] == 1:
    #         Y[i] = 0
    #     else:
    #         Y[i] = 1

    ##### ----- #####

    return X, Y, beta

if __name__ == "__main__":
    theta = 0.3
    n = 10
    m = 5
    X, Y_true, beta_true = logisticregr(theta, n, m)
    print("shape of X is ", X.shape)
    print("shape of beta is ", beta_true.shape)
    print("shape of Y is ", Y_true.shape)
    print("X matrix is \n", X)
    print("beta vector is \n", beta_true)
    print("Y dependent vector is \n ", Y_true)

```

2 and 4)# Question2 and Question 4.py

```

import argparse # Commandline input

import numpy as np
from numpy.linalg import norm

```

```

import Q1 as q1 # importing First Question solution

parser = argparse.ArgumentParser()
parser.add_argument(
    "-regularization",
    type=str,
    default="0",
    help="Put 11 for lasso or 12 for Ridge or 0 for no regularization",
)
parser.add_argument(
    "-regconstant",
    type=float,
    help="Regularization constant (float)",
)
args = parser.parse_args()

def cost(Y_pred, Y_true, beta):
    """Calculate the cost

    Args:
        Y_pred (array): predicted value of Y output
        beta (array): value of beta array at current iteration including
initial initialization

    Returns:
        float: cost at current iteration
    """

    # remove 0 and 1 from Y_pred and convert it to close to 0 and 1
respectively
    Y_pred = np.where(Y_pred == 0, 0.01, Y_pred)
    Y_pred = np.where(Y_pred == 1, 0.99, Y_pred)

    # removing any 0 component in beta such that 0 doesn't come in
denominator in regularization term
    beta = np.where(beta == 0, 0.01, beta)

    if args.regularization == "0":

```

```

        cost = -np.mean(Y_true * np.log(Y_pred) + (1 - Y_true) * (np.log(1
- Y_pred))) # initial cost function

    elif args.regularization == "l1":
        cost = -np.mean(Y_true * np.log(Y_pred) + (1 - Y_true) * (np.log(1
- Y_pred))) + args.regconstant * np.sum(
            np.abs(beta)
        ) # initial cost function
    elif args.regularization == "l2":
        cost = -np.mean(Y_true * np.log(Y_pred) + (1 - Y_true) * (np.log(1
- Y_pred))) + args.regconstant * np.sum(
            beta**2
        )

    return cost

def gradientdescent(X, Y_true, epochs, threshold, LR):
    np.random.seed(1998)
    beta = np.random.rand(X.shape[1], 1) # initial values of beta
coefficients
    # print("Initial beta ", beta)

    Y_pred = 1 / (1 + np.exp(-np.dot(X, beta))) # initial Y prediction

    init_cost = cost(Y_pred, Y_true, beta) # initial cost function

    # print("init cost", init_cost)
    for j in range(epochs):
        if args.regularization == "0":
            beta[0] = beta[0] + LR * sum(Y_true - Y_pred) # intercept

            for i in range(1, X.shape[1]):
                residualderivative = -(Y_true - Y_pred) * np.array(X[:,
i]).reshape(Y_true.shape[0], 1)
                beta[i] = beta[i] - LR * sum(residualderivative)

            if args.regularization == "l1":
                beta[0] = beta[0] + LR * sum(Y_true - Y_pred) -
args.regconstant * beta[0] / np.abs(beta[0]) # intercept

```

```

        for i in range(1, X.shape[1]):
            residualderivative = -(Y_true - Y_pred) * np.array(X[:,
i]).reshape(
                Y_true.shape[0], 1
            ) + args.regconstant * beta[i] / np.abs(beta[i])
            beta[i] = beta[i] - LR * sum(residualderivative)

    if args.regularization == "l2":
        beta[0] = beta[0] + LR * sum(Y_true - Y_pred) -
args.regconstant * 2 * beta[0] # intercept

        for i in range(1, X.shape[1]):
            residualderivative = (
                -2 * (Y_true - Y_pred) * np.array(X[:,
i]).reshape(Y_true.shape[0], 1)
                + args.regconstant * 2 * beta[i]
            )
            beta[i] = beta[i] - LR * sum(residualderivative)

    Y_pred = 1 / (1 + np.exp(-np.dot(X, beta))) # Y prediction for the
new betas

    new_cost = cost(Y_pred, Y_true, beta) # Cost function for the new
betas

    # print(new_cost)
    if init_cost - new_cost > threshold: # condition
        init_cost = new_cost
        if j == epochs - 1:
            print(f"Loop ran till last epoch {j}")
        else:
            print(f"Loop break at {j}th iteration")
            break

    # cosinesim_beta = np.dot(beta_true.flatten(), beta.flatten()) /
(norm(beta_true.flatten()) * norm(beta.flatten()))
    Y_pred = Y_pred >= 0.5
    Y_pred = Y_pred.astype(int)

    cosinesim_Y = np.dot(Y_true.flatten(), Y_pred.flatten()) /
(norm(Y_true.flatten()) * norm(Y_pred.flatten()))

```

```


    return beta, new_cost, j, cosinesim_Y

if __name__ == "__main__":

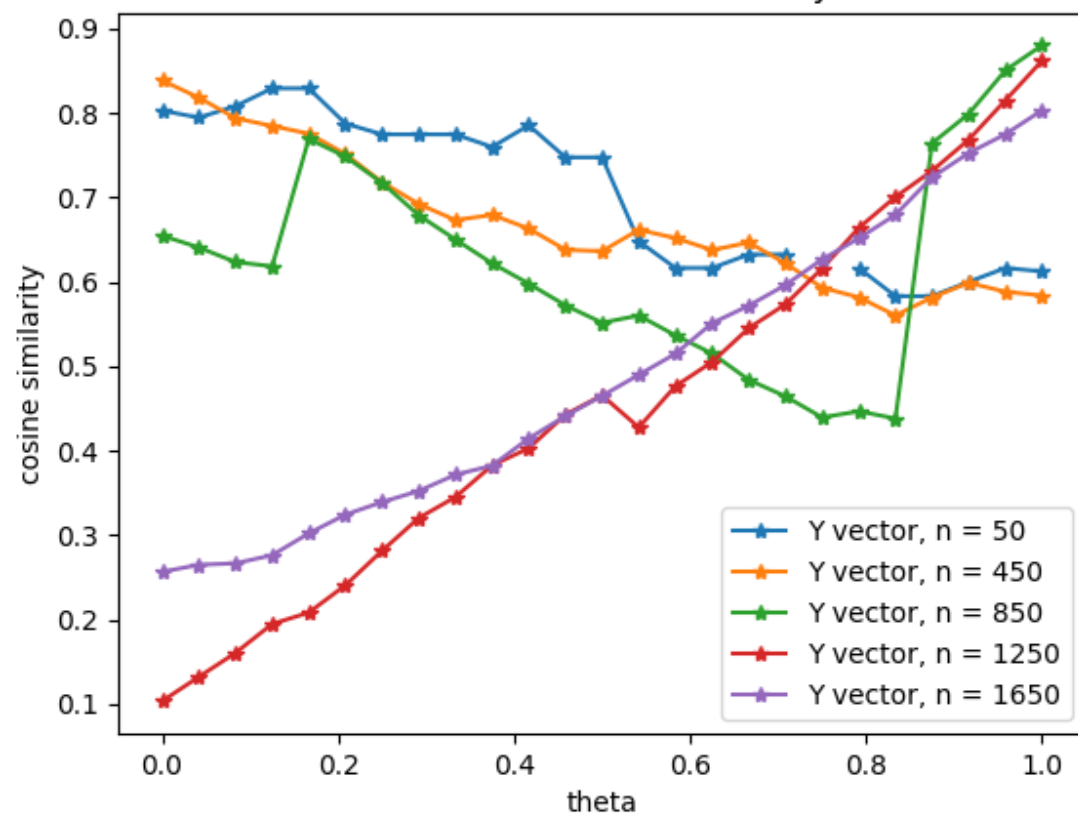
    theta = 0.8
    n = 1000
    m = 50
    X, Y_true, beta_true = q1.logisticregr(theta, n, m)
    print(f"True (Original) beta vector: {beta_true}")
    print(f"True (Original) output vector: {Y_true}")
    beta_new, new_cost, iteration, cosinesim_Y = gradientdescent(X, Y_true,
epochs=10000, threshold=0.00001, LR=0.001)
    print(f"Obtained beta vector after gradient descent: {beta_new}")
    print(f"cost function value is {new_cost}")
    # print(f"Cosine similarity between True beta vector and beta vector
after gradient descent is {cosinesim_beta}")
    print(f"Cosine similarity between True Y vector and Y vector after
gradient descent is {cosinesim_Y}")


```

Variation of cosine similarity between Y_{true} and Y_{pred} considering L1 regularization:




Variation of cosine similarity





Variation of Cosine similarity using L2 regularization:



3)

```
# Question3.py
import matplotlib.pyplot as plt
import numpy as np

import Q1 as q1 # importing First Question solution
import Q2_Q4 as q2

# -----Caveat----- #

# 1) I have divided the script into three parts, please run one by one and
comment others which is not needed for the moment
# 2) section 1 and section 2 took 5-10 mins to run when theta but this may
vary depending on n and theta values you put
# 3) section 3 will take the most amount of time, on my system it took
around ~25 mins again depending on the parameters you put
```

```

# ----- #

if __name__ == "__main__":

    # -----#

    #
    #          VARIATION WITH theta
    #
    # -----#

    # m = 5
    # iter = [] # list to store at which epoch , we break the loop to get
    # final new beta
    # costf = [] # final cost functions list

    # for n in range(50, 2000, 400):
    #     betacos_similarity = []
    #     Ycos_similarity = []
    #     for theta in np.linspace(0, 1, 25): # Variation with theta,
    #         0<=theta<=1
    #             print(f"For n {n} and theta {theta}")
    #             X, Y_true, beta_true = q1.logisticregr(theta, n, m)
    #             beta_new, new_cost, iteration, cosinesim_Y =
    #             q2.gradientdescent(X, Y_true, 10000, 0.000001, 0.001)
    #             Ycos_similarity.append(cosinesim_Y)

    #     plt.plot(np.linspace(0, 1, 25), Ycos_similarity, marker="*",
    # label="Y vector, n = " + str(n))
    #     plt.ylabel("cosine similarity")
    #     plt.xlabel("theta")
    #     plt.title("Variation of cosine similarity")
    #     plt.legend()
    #     plt.show()

    # -----#

    #
    #          VARIATION WITH N
    #
    # -----#

```

```

# m = 5

# for theta in np.linspace(0, 1, 4):
#     betacos_similarity = []
#     Ycos_similarity = []
#     for n in range(100, 2000, 400): # Variation with n
#         print(f"For theta {theta} and n {n}")
#         X, Y_true, beta_true = q1.logisticregr(theta, n, m)
#         beta_new, new_cost, iteration, cosinesim_Y =
q2.gradientdescent(X, Y_true, 10000, 0.000001, 0.001)
#         # betacos_similarity.append(cosinesim_beta)
#         Ycos_similarity.append(cosinesim_Y)

#     # plt.plot(range(10, 2000, 100), betacos_similarity, marker="*",
label="Beta vector, theta = " + str(theta))
#     plt.plot(range(100, 2000, 400), Ycos_similarity, marker="*",
label="Y vector, theta = " + str(round(theta, 2)))
# plt.xlabel("n")
# plt.ylabel("cosine similarity")
# plt.title("Variation of cosine similarity")
# plt.legend()
# plt.show()

# -----#

#         VARIATION WITH BOTH theta and N

# -----#

ninputs = np.concatenate((np.linspace(10, 1999, 10, dtype=int),
np.linspace(2000, 20000, 10, dtype=int))) # non-linear spacing
theta = np.linspace(0.01, 0.99, 10)
cosinesim_Y = np.zeros((10, 20))
cosinesim_beta = np.zeros((10, 20))
m = 10
for t in theta:
    for n in ninputs:
        print(f"For n {n} and theta {t}")
        X, Y_true, beta_true = q1.logisticregr(t, n, m)
        (

```

```

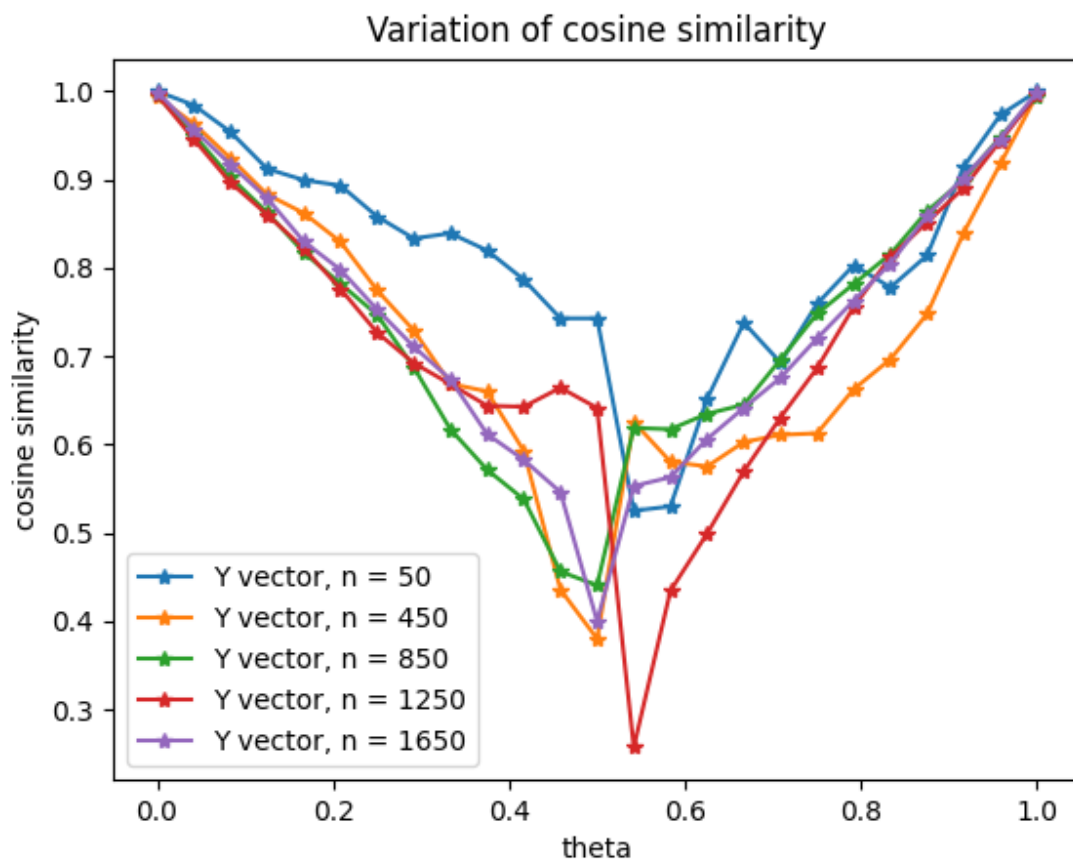
        beta_new,
        new_cost,
        iteration,
        cosinesim_Y[np.where(theta == t)[0], np.where(ninputs ==
n)[0]],


    ) = q2.gradientdescent(X, Y_true, 10000, 0.000001, 0.001)

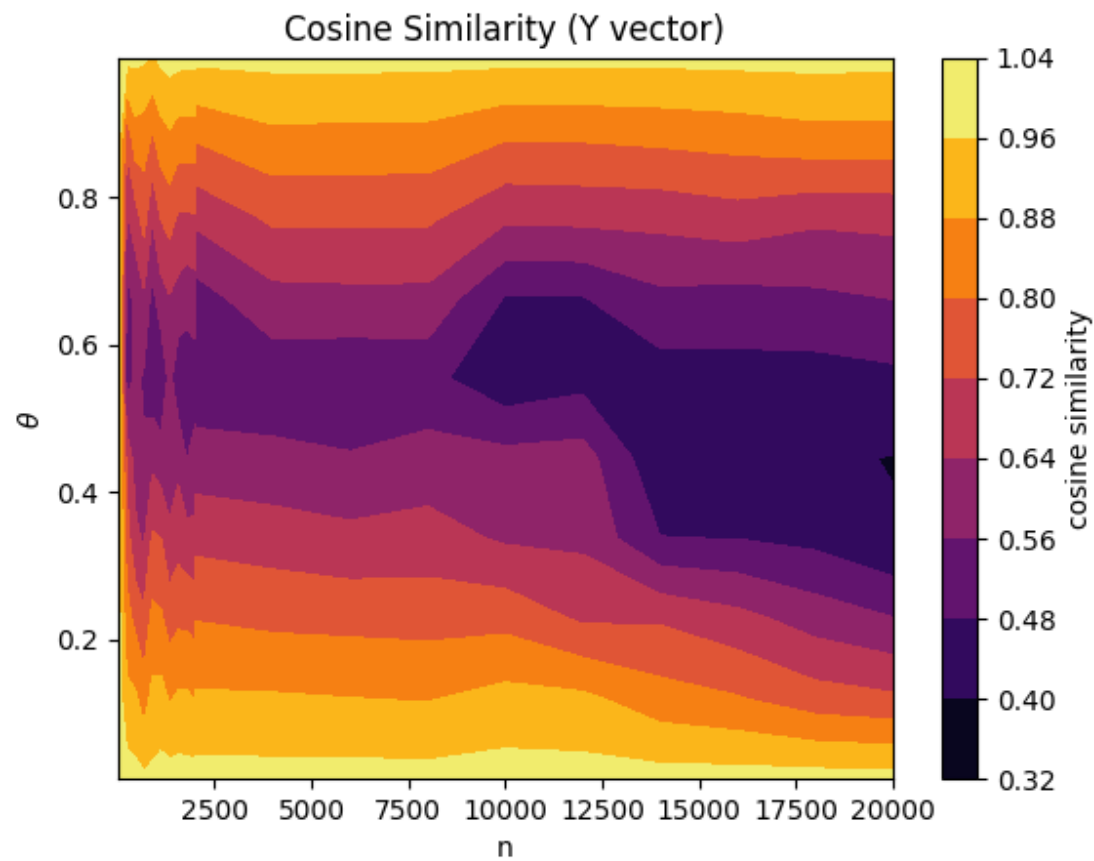
plt.title("Cosine Similarity (Y vector)")
plt.contourf(ninputs, theta, cosinesim_Y, cmap="inferno")
c = plt.colorbar()
c.set_label("cosine similarity")
plt.xlabel("n")
plt.ylabel(r"$\theta$")
plt.show()

```

Cosine similarity between Y_{true} and $Y_{\text{prediction}}$:







5)

```
# Question5.py
import argparse # Commandline input

import numpy as np
from numpy.linalg import norm

parser = argparse.ArgumentParser()
parser.add_argument(
    "-regression",
    type=str,
    default="linear",
```



```

        help="Put linear for linear regression and logistic for logistic
regression",
    )
    parser.add_argument(
        "-regularization",
        type=str,
        default="0",
        help="Put 11 for lasso or 12 for Ridge or 0 for no regularization",
    )
    parser.add_argument(
        "-regconstant",
        type=float,
        help="Regularization constant (float)",
    )
    args = parser.parse_args()

class linear_logistic_reg:
    def __init__(self, std_dev, theta, n, m, epochs, threshold, LR):
        self.std_dev = std_dev
        self.theta = theta
        self.n = n
        self.m = m
        self.epochs = epochs
        self.threshold = threshold
        self.LR = LR

    def Regression(self):
        """Function to generate m*n independent variables, n*1 dependent
variable, and coefficient matrix

        Args:
            stddev (float): spread of noise in the output variable (linear
regression)
            theta (float): spread of noise in the output variable
(logisitic regression)
            n (int): size of the data set
            m (int): number of indepedent variables

        Returns:

```

```

        array: independent, dependent variables matrix and coefficient
matrix
        """
        x_ones = np.ones((self.n, 1)) # first column of ones
        # np.random.seed(1998)
        x_random = np.random.randn(self.n, self.m) # (n*m) matrix of
random Real numbers
        X = np.concatenate((x_ones, x_random), axis=1) # n rows and m+1
columns where 1st column is of ones
        beta = np.random.rand(self.m + 1, 1) # Coefficients

        if args.regression == "linear":
            error = np.random.normal(loc=0, scale=self.std_dev,
size=(self.n, 1))
            Y = (
                np.matmul(X, beta) + error
            ) # X*beta i.e. matrix multiplication of (n*m) and (m * 1)
shape= Y shape (n*1)

        if args.regression == "logistic":
            Y = (
                1 / (1 + np.exp(-np.dot(X, beta))) >= 0.5
            ) # X*beta i.e. matrix multiplication of (n*m) and (m * 1)
shape= Y shape (n*1)

            Y = Y.astype(int)
            noise = np.random.binomial(n=1, p=self.theta, size=(self.n, 1))
            Y = Y + noise
            Y %= 2

        return X, Y, beta

def cost(self, Y_pred, Y_true, beta):
    """Calculate the cost

    Args:
        Y_pred (array): predicted value of Y output
        beta (array): value of beta array at current iteration
including initial initialization

```

```

Returns:
    float: cost at current iteration
    """

    # remove 0 and 1 from Y_pred and convert it to close to 0 and 1
    respectively
    Y_pred = np.where(Y_pred == 0, 1e-4, Y_pred)
    Y_pred = np.where(Y_pred == 1, 1 - 1e-4, Y_pred)

    # removing any 0 component in beta such that 0 doesn't come in
    denominator in regularization term
    beta = np.where(beta == 0, 0.01, beta)

    if args.regularization == "0":
        cost = -np.mean(Y_true * np.log(Y_pred) + (1 - Y_true) *
            (np.log(1 - Y_pred))) # initial cost function
    elif args.regularization == "l1":
        cost = -np.mean(Y_true * np.log(Y_pred) + (1 - Y_true) *
            (np.log(1 - Y_pred))) + args.regconstant * np.sum(
            np.abs(beta)
        ) # initial cost function
    elif args.regularization == "l2":
        cost = -np.mean(Y_true * np.log(Y_pred) + (1 - Y_true) *
            (np.log(1 - Y_pred))) + args.regconstant * np.sum(
            beta**2
        )

    return cost

def gradientdescent(self, X, Y_true):
    np.random.seed(1998)
    beta = np.random.rand(X.shape[1], 1) # initial values of beta
    coefficients

    print("Initial beta ", beta)

    if args.regression == "linear":
        Ypredict = np.matmul(X, beta) # initial Y prediction
        init_cost = sum((Y_true - Ypredict) ** 2) # initial cost
function

```

```

        print("init cost", init_cost)
        for j in range(self.epochs):
            beta[0] = beta[0] + 2 * self.LR * sum(Y_true - np.matmul(X,
beta)) # intercept
            for i in range(1, X.shape[1]):
                residualderivative = (
                    -2 * (Y_true - np.matmul(X, beta)) * np.array(X[:,
i]).reshape(Y_true.shape[0], 1)
                )
                beta[i] = beta[i] - self.LR * sum(residualderivative)

            Ypred = np.matmul(X, beta) # Y prediction or the new betas
            new_cost = sum((Y_true - Ypred) ** 2) # Cost function for
the new betas

            print(new_cost)
            if init_cost - new_cost > self.threshold: # condition
                init_cost = new_cost
            else:
                print(f"Loop break at {j}th iteration")
                break

            cosinesim_Y = np.dot(Y_true.flatten(), Ypred.flatten()) /
(norm(Y_true.flatten()) * norm(Ypred.flatten()))

            return beta, new_cost, j, cosinesim_Y

    if args.regression == "logistic":

        Y_pred = 1 / (1 + np.exp(-np.dot(X, beta))) # initial Y
prediction

        init_cost = test.cost(Y_pred, Y_true, beta) # initial cost
function

        # print("init cost", init_cost)
        for j in range(self.epochs):
            if args.regularization == "0":
                beta[0] = beta[0] + self.LR * sum(Y_true - Y_pred) #
intercept

```

```

        for i in range(1, X.shape[1]):
            residualderivative = -(Y_true - Y_pred) *
np.array(X[:, i]).reshape(Y_true.shape[0], 1)
            beta[i] = beta[i] - self.LR *
sum(residualderivative)

        if args.regularization == "l1":
            beta[0] = (
                beta[0] + self.LR * sum(Y_true - Y_pred) -
args.regconstant * beta[0] / np.abs(beta[0])
            ) # intercept

            for i in range(1, X.shape[1]):
                residualderivative = -(Y_true - Y_pred) *
np.array(X[:, i]).reshape(
                    Y_true.shape[0], 1
                ) + args.regconstant * beta[i] / np.abs(beta[i])
                beta[i] = beta[i] - self.LR *
sum(residualderivative)

            if args.regularization == "l2":
                beta[0] = beta[0] + self.LR * sum(Y_true - Y_pred) -
args.regconstant * 2 * beta[0] # intercept

                for i in range(1, X.shape[1]):
                    residualderivative = (
                        -2 * (Y_true - Y_pred) * np.array(X[:,
i]).reshape(Y_true.shape[0], 1)
                        + args.regconstant * 2 * beta[i]
                    )
                    beta[i] = beta[i] - self.LR *
sum(residualderivative)

            Y_pred = 1 / (1 + np.exp(-np.dot(X, beta))) # Y prediction
for the new betas

            new_cost = test.cost(Y_pred, Y_true, beta) # Cost function
for the new betas

            # print(new_cost)
            if init_cost - new_cost > self.threshold: # condition

```

```

        init_cost = new_cost
        if j == self.epochs - 1:
            print(f"Loop ran till last epoch {j}")
        else:
            print(f"Loop break at {j}th iteration")
            break

        # cosinesim_beta = np.dot(beta_true.flatten(), beta.flatten())
        / (norm(beta_true.flatten()) * norm(beta.flatten()))
        Y_pred = Y_pred >= 0.5
        Y_pred = Y_pred.astype(int)

        cosinesim_Y = np.dot(Y_true.flatten(), Y_pred.flatten()) /
        (norm(Y_true.flatten()) * norm(Y_pred.flatten()))
        return beta, new_cost, j, cosinesim_Y

if __name__ == "__main__":

    test = linear_logistic_reg(std_dev=0.3, theta=0.8, n=1000, m=50,
epochs=10000, threshold=0.00001, LR=0.001)
    X, Y_true, beta_true = test.Regression()
    print(f"True (Original) beta vector: {beta_true}")
    print(f"True (Original) output vector: {Y_true}")
    beta_new, new_cost, iteration, cosinesim_Y = test.gradientdescent(
        X, Y_true
    )
    print(f"Obtained beta vector after gradient descent: {beta_new}")
    print(f"cost function value is {new_cost}")
    print(f"Cosine similarity between True Y vector and Y vector after
gradient descent is {cosinesim_Y}")

```

$$\textcircled{3} \text{ Loss function}_{(L)} = -[y \log \hat{y} + (1-y) \log (1-\hat{y})]$$

$y \rightarrow \text{data}$

$$\hat{y} = \sigma(z) = \frac{1}{1+e^{-z}} \quad \text{where } z = \vec{x} \cdot \vec{\beta} = \sum_i \beta_i x_i$$

$$\frac{\partial L}{\partial \beta_i} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z} \frac{\partial z}{\partial \beta_i}$$

$$\frac{\partial L}{\partial \hat{y}} = \frac{\partial}{\partial \hat{y}} (y \log \hat{y} + (1-y) \log (1-\hat{y}))$$

$$= - \left[\frac{y}{\hat{y}} + \frac{(1-y)}{1-\hat{y}} (-1) \right]$$

$$= - \frac{y}{\hat{y}} + \frac{1-y}{1-\hat{y}}$$

$$\frac{\partial \hat{y}}{\partial z} = \frac{1/(1+e^{-z})}{\partial z} = - (1+e^{-z})^2 \frac{\partial (1+e^{-z})}{\partial z}$$

$$= - (1+e^{-z})^2 e^{-z} \frac{\partial (-z)}{\partial z}$$

$$= - (1+e^{-z})^2 e^{-z} = \frac{1}{1+e^{-z}} \left(\frac{e^{-z}}{1+e^{-z}} \right)$$

$$= \hat{y} (1-\hat{y})$$

$$\therefore \frac{\partial L}{\partial \beta_i} = \left(-\frac{y}{\hat{y}} + \frac{1-y}{1-\hat{y}} \right) (\hat{y}(1-\hat{y})) \frac{\partial z}{\partial \beta_i}$$

AS

$$\text{AS } x_0 = 1 \quad \therefore \frac{\partial z}{\partial \beta_0} = \frac{\partial \beta_0}{\partial \beta_0} = 1$$

\forall other x_n , where $n \neq 0$

$$\frac{\partial z}{\partial \beta_i} = x_i$$

$$\therefore \frac{\partial L}{\partial \beta_i} = \begin{cases} (\hat{y} - y) & \text{where } i = 0 \\ (\hat{y} - y) x_i; & \forall i \neq 0 \end{cases}$$

(ii) Cosine Similarity $\left(\frac{\vec{A} \cdot \vec{B}}{|\vec{A}| |\vec{B}|} = \cos \theta \right)$; where θ is the Angle b/w 2 vectors, is minima at $\theta = 0.5$

here $\vec{A} = Y_{\text{true}}$

$\vec{B} = Y_{\text{predicted}}$